

Predicting Customer Churn

Identifying Customer that are susceptible to Churn(leave) in order to enhance retention strategies and boost Sales/Business Growth

Import packages and libraries

```
In [1]: from ast import literal_eval #literal_eval allows computation on strings
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt #to visualize
import seaborn as sns
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier #helps predict
from tqdm import tqdm #helps visualize the progression of loops and action in py

import warnings      #to avoid seeing warings in pandas packages
warnings.filterwarnings('ignore')
```

```
In [2]: #data collection
#Load the dataset into Python with pandas
df = pd.read_excel("Dataset.xlsx")
```

Data Exploration

In this section:

- understanding the data structure, statictics and quality of the dataset
- visualizing the data to gain insughts
- checking for missing values

```
In [3]: #show the fisrt 5 rows of the dataset
df.head()
```

Out[3]:

	CustomerID	Name	Age	Gender	Location	Email
0	1001	Mark Barrett	31	Male	Andrewfort	allison74@example.net
1	1002	Jeremy Welch	66	Female	Millerhaven	fmiller@example.com
2	1003	Brandon Patel	36	Female	Lozanostad	jasonbrown@example.org
3	1004	Tina Martin	62	Female	South Dustin	matthew62@example.net
4	1005	Christopher Rodriguez	68	Female	West James	shannonstrickland@example.org

5 rows × 21 columns

In [4]:

```
#handling missing values
missing_values = df.isnull().sum()
missing_values
```

```
Out[4]: CustomerID      0  
Name          0  
Age          0  
Gender        0  
Location       0  
Email          0  
Phone          0  
Address         0  
Segment         0  
PurchaseHistory  0  
SubscriptionDetails 0  
ServiceInteractions 0  
PaymentHistory    0  
WebsiteUsage      0  
ClickstreamData     0  
EngagementMetrics   0  
Feedback          0  
MarketingCommunication 0  
NPS              0  
ChurnLabel        0  
Timestamp         0  
dtype: int64
```

```
In [5]: #check for duplicate values  
df.duplicated().any()
```

```
Out[5]: False
```

```
In [6]: #statistics overview of the numerical values of the dataset  
df.describe()
```

```
Out[6]:
```

	CustomerID	Age	NPS	ChurnLabel
count	12483.00000	12483.00000	12483.00000	12483.00000
mean	7242.00000	43.930065	2.973884	0.505808
std	3603.67604	15.341521	2.644623	0.499986
min	1001.00000	18.000000	0.000000	0.000000
25%	4121.50000	31.000000	1.000000	0.000000
50%	7242.00000	44.000000	2.000000	1.000000
75%	10362.50000	57.000000	4.000000	1.000000
max	13483.00000	70.000000	9.000000	1.000000

How the target variable, CHURN LABEL are distributed

```
In [7]: #setup the figure and axes  
fig, ax = plt.subplots(nrows=2, ncols=2, figsize=(15, 10))  
  
#Plot the distribution of the target variable ChurnLabel  
sns.countplot(x='ChurnLabel', data=df, ax=ax[0,0], palette=['green', 'pink'])  
ax[0, 0].set_title('Distribution of ChurnLabel')  
ax[0, 0].set_xticklabels(['No Churn', 'Churn'])
```

```

#Plot distribution of the Gender
sns.countplot(x='Gender', data=df, ax=ax[0, 1], palette=['blue', 'orange'])
ax[0, 1].set_title('Distribution of Gender')


#Plot distribution of Segment
sns.countplot(x='Segment', data=df, ax=ax[1, 0], palette='Set3')
ax[1, 0].set_title('Distribution of Segment')

#Age distribution
sns.histplot(df['Age'], bins=30, ax=ax[1, 1], kde=True, color='red')
ax[1, 1].set_title('Age Distribution')

plt.tight_layout()
plt.show

```

Out[7]: <function matplotlib.pyplot.show(close=None, block=None)>



Correlation Analysis: Which columns in the dataset correlates with the 'ChurnLabel' column.


```

In [8]: #Calculate the correlation with the ChurnLabel
correlation = df.select_dtypes(include=np.number).corr()['ChurnLabel'].sort_values(ascending=False)

#Plot the correlation
plt.figure(figsize=(10, 6))
correlation.drop('ChurnLabel').plot(kind='bar', color='skyblue')
plt.title('Correlation with ChurnLabel')
plt.ylabel('Correlation Coefficient')
plt.show

```

Out[8]: <function matplotlib.pyplot.show(close=None, block=None)>



I am exploring the nested columns and see how they are nested to further explore the dataset

```
In [9]: #List out the nested columns
nested_columns = [
    'PurchaseHistory',
    'SubscriptionDetails',
    'ServiceInteractions',
    'PaymentHistory',
    'WebsiteUsage',
    'ClickstreamData',
    'EngagementMetrics',
    'Feedback',
    'MarketingCommunication'
]

#print the first value in each of the columns for a detailed view
w1, w2 = 25, 1000 #(width , name)
for col in nested_columns:
    row = [col, df[col][0]]
    print ('\n| {:<{w1}} | {:<{w2}} | '.format(*row, w1=w1, w2=w2))
```

```
| PurchaseHistory           | [{"Product": "Frozen Cocktail Mixes", "Frequency": 8, "Value": 884.43}, {"Product": "Printer, Copier & Fax Machine Accessories", "Frequency": 7, "Value": 397.14}, {"Product": "Hockey Stick Care", "Frequency": 10, "Value": 498.92}, {"Product": "Guacamole", "Frequency": 2, "Value": 718.43}, {"Product": "Mortisers", "Frequency": 2, "Value": 614.08}, {"Product": "Rulers", "Frequency": 6, "Value": 221.68}, {"Product": "Invitations", "Frequency": 3, "Value": 660.04}]  
|  
| SubscriptionDetails      | {"Plan": "Express", "Start_Date": "2020-06-08", "End_Date": "2022-10-27"}  
|  
| ServiceInteractions       | [{"Type": "Call", "Date": "2019-09-26"}, {"Type": "Chat", "Date": "2021-07-25"}, {"Type": "Email", "Date": "2020-04-13"}, {"Type": "Chat", "Date": "2020-11-15"}]  
|  
| PaymentHistory            | [{"Method": "Credit Card", "Late_Payments": 5}, {"Method": "PayPal", "Late_Payments": 11}, {"Method": "Bank Transfer", "Late_Payments": 24}]  
|  
| WebsiteUsage              | {"PageViews": 49, "TimeSpent(minutes)": 15}  
|  
| ClickstreamData           | [{"Action": "Add to Cart", "Page": "register", "Timestamp": "2020-09-13 17:06:44"}, {"Action": "Search", "Page": "login", "Timestamp": "2022-03-30 14:51:52"}, {"Action": "Click", "Page": "about", "Timestamp": "2019-11-10 05:48:48"}, {"Action": "Add to Cart", "Page": "terms", "Timestamp": "2019-05-15 10:17:44"}, {"Action": "Add to Cart", "Page": "author", "Timestamp": "2022-07-14 03:40:53"}, {"Action": "Search", "Page": "main", "Timestamp": "2019-01-13 08:39:42"}, {"Action": "Add to Cart", "Page": "faq", "Timestamp": "2019-02-19 05:28:25"}, {"Action": "Add to Cart", "Page": "about", "Timestamp": "2020-11-01 20:59:55"}, {"Action": "Click", "Page": "faq", "Timestamp": "2021-12-22 16:39:40"}, {"Action": "Add to Cart", "Page": "main", "Timestamp": "2020-11-11 03:25:36"}, {"Action": "Click", "Page": "privacy", "Timestamp": "2021-06-13 06:18:41"}, {"Action": "Add to Cart", "Page": "search", "Timestamp": "2022-03-28 16:25:35"}, {"Action": "Search", "Page": "homepage", "Timestamp": "2019-09-26 12:27:42"}, {"Action": "Click", "Page": "search", "Timestamp": "2021-03-31 16:35:39"}, {"Action": "Search", "Page": "main", "Timestamp": "2021-12-22 10:02:19"}, {"Action": "Search", "Page": "about", "Timestamp": "2019-08-24 05:11:40"}, {"Action": "Add to Cart", "Page": "index", "Timestamp": "2021-04-30 00:38:03"}, {"Action": "Search", "Page": "privacy", "Timestamp": "2021-06-21 16:23:49"}, {"Action": "Search", "Page": "about", "Timestamp": "2022-04-03 07:25:20"}, {"Action": "Search", "Page": "author", "Timestamp": "2022-11-07 02:24:31"}, {"Action": "Search", "Page": "about", "Timestamp": "2019-08-25 17:37:59"}, {"Action": "Search", "Page": "post", "Timestamp": "2020-12-18 01:36:34"}, {"Action": "Search", "Page": "home", "Timestamp": "2021-11-24 07:33:26"}, {"Action": "Search", "Page": "login", "Timestamp": "2020-11-15 07:21:21"}]  
|  
| EngagementMetrics          | {"Logins": 19, "Frequency": "Weekly"}  
|  
| Feedback                   | {"Rating": 1, "Comment": "I move baby go small big. Office institution six. Fact until hear technology right company seek."}  
|  
| MarketingCommunication     | [{"Email_Sent": "2019-10-17", "Email_Opened": "2022-01-12", "Email_Clicked": "2022-11-27"}, {"Email_Sent": "2019-10-17", "Email_Opened": "2022-01-12", "Email_Clicked": "2022-11-27"}]
```

```

ed': '2022-01-12', 'Email_Clicked': '2022-11-27'}, {'Email_Sent': '2019-10-17',
'Email_Opened': '2022-01-12', 'Email_Clicked': '2022-11-27'}, {'Email_Sent': '2019-10-17', 'Email_Opened': '2022-01-12', 'Email_Clicked': '2022-11-27'},
{'Email_Sent': '2019-10-17', 'Email_Opened': '2022-01-12', 'Email_Clicked': '2022-11-27'}, {'Email_Sent': '2019-10-17', 'Email_Opened': '2022-01-12', 'Email_Clicked': '2022-11-27'},
{'Email_Sent': '2019-10-17', 'Email_Opened': '2022-01-12', 'Email_Clicked': '2022-11-27'}]
|

```

Temporal Analysis: I looked at how the Churn rate chances over time to see if there are any reoccuring patterns

```


In [10]: #Convert Timestamp to datetime format
df['Timestamp'] = pd.to_datetime(df['Timestamp'])

#Extract the month-year from the Timestamp
df['MonthYear'] = df['Timestamp'].dt.to_period('M')

# Group by MonthYear and calculate the churn rate
monthly_churn_rate = df.groupby('MonthYear')[['ChurnLabel']].sum()

# Plot the Churn rate over time
plt.figure(figsize=(14,7))
monthly_churn_rate.plot()
plt.title('Monthly Churn Rate Over Time')
plt.ylabel('Number of Churned Customers')
plt.xlabel('Month-Year')
plt.show()

```



I need to see if there are any patterns between the customers feedback (Rating) and the ChurnLabel

```

In [11]: # Extract the ratings from the feedback column and create a new column for it
df['FeedbackRating'] = df['Feedback'].apply(lambda x: eval(x)[ 'Rating']) #eval

# Plot the relationship btwn feedback rating and ChurnLabel
plt.figure(figsize=(12, 6))
sns.countplot(x='FeedbackRating', data=df, hue='ChurnLabel', palette='Set2')


```

```

plt.title('Relationship between Feedback Ratings and ChurnLabel')
plt.xlabel('Rating')
plt.ylabel('Count')
plt.legend(title = 'ChurnLabel', loc='upper right', labels=['No Churn', 'Churn'])
plt.show

```

Out[11]: <function matplotlib.pyplot.show(close=None, block=None)>



The visual shows there isn't much indication that feedback rating affects the Churn rate based on the positive or negative ratings

Data Preprocessing & Feature Engineering

Here I will be:

- create new feature that may have predictive power,
- convert categorical variable to numeric variables, using encoding techniques,
- scaling or normalize numeric variables where necessary,
- split data into train & test subsets
- remove irrelevant features.

I will start by converting the nested values from string formats to lists/dictionaries using the literal_eval function

In [12]: #Listing out the nested columns
nested_columns = [
 'PurchaseHistory',
 'SubscriptionDetails',
 'ServiceInteractions',
 'PaymentHistory',
 'WebsiteUsage',
 'ClickstreamData',
 'EngagementMetrics',
 'Feedback',
 'MarketingCommunication'
]

```
#Applying the literal_eval function to convert to list/dict
for feature in nested_columns:
    df[feature] = df[feature].apply(literal_eval)
```

More features to be extracted from the dataset

```
In [13]: # PurchaseHistory extracts all product purchased
df['PurchasedProducts'] = df['PurchaseHistory'].apply(lambda x: '|'.join([i['Product'] for i in x]))
df['PurchaseFrequency'] = df['PurchaseHistory'].apply(lambda x: sum([i['Frequency'] for i in x]))
df['PurchaseValue'] = df['PurchaseHistory'].apply(lambda x: sum([i['Value'] for i in x]))

#SubscriptionDetails
df['SubscriptionPlan'] = df['SubscriptionDetails'].apply(lambda x: x['Plan'])
df['SubscriptionStartDate'] = df['SubscriptionDetails'].apply(lambda x: x['Start Date'])
df['SubscriptionEndDate'] = df['SubscriptionDetails'].apply(lambda x: x['End Date'])
df['SubscriptionDuration'] = (pd.to_datetime(df['SubscriptionEndDate']) - pd.to_datetime(df['SubscriptionStartDate']))

#WebsiteUsage
df['WebsitePageViews'] = df['WebsiteUsage'].apply(lambda x: x['PageViews'])
df['WebsiteTimeSpent'] = df['WebsiteUsage'].apply(lambda x: x['TimeSpent'](minutes))

#EngagementMetrics
df['EngagementMetricsLogins'] = df['EngagementMetrics'].apply(lambda x: x['Logins'])
df['EngagementMetricsFrequency'] = df['EngagementMetrics'].apply(lambda x: x['Frequency'])

#Feedbacks
df['FeedbackRating'] = df['Feedback'].apply(lambda x: x['Rating'])
df['FeedbackComment'] = df['Feedback'].apply(lambda x: x['Comment'])

#MarketingCommunication
df['MarketingCommunicationNoOfEmails'] = df['MarketingCommunication'].apply(lambda x: len(x))
df['MarketingCommunicationOpenClickDiff'] = df['MarketingCommunication'].apply(lambda x: np.mean([
    (pd.to_datetime(i['Email_Clicked']) - pd.to_datetime(i['Email_Opened'])).days for i in x
]))
df['MarketingCommunicationSentOpenDiff'] = df['MarketingCommunication'].apply(lambda x: np.mean([
    (pd.to_datetime(i['Email_Opened']) - pd.to_datetime(i['Email_Sent'])).days for i in x
]))
```

Special Extraction from three columns:

- ServiceInteractions,
- PaymentHistory,
- ClickstreamData.

Check for the unique parameters/values in these columns.

```
In [14]: # Get all unique ServiceInteraction 'Types'
service_interaction_types = df['ServiceInteractions'].apply(lambda x: list(set([
    service_interaction_type for service_interaction_type in x])))
```

```

unique_service_interaction_type = []
for i in service_interaction_types:
    unique_service_interaction_type.extend(i)
unique_service_interaction_type = list(set(unique_service_interaction_type))
print('All unique Service Interaction Types:', unique_service_interaction_type)

#Get all unique PaymentHistory 'Method'
payment_history_methods = df['PaymentHistory'].apply(lambda x: list(set([i['Method'] for i in x])))
payment_history_methods = payment_history_methods.to_list()
unique_payment_history_methods = []
for i in payment_history_methods:
    unique_payment_history_methods.extend(i)
unique_payment_history_methods = list(set(unique_payment_history_methods))
print('All unique Payment History Methods:', unique_payment_history_methods)

#Get all unique ClickstreamData 'Action'
clickstream_data_actions = df['ClickstreamData'].apply(lambda x: list(set([i['Action'] for i in x])))
clickstream_data_actions = clickstream_data_actions.to_list()
unique_clickstream_data_actions = []
for i in clickstream_data_actions:
    unique_clickstream_data_actions.extend(i)
unique_clickstream_data_actions = list(set(unique_clickstream_data_actions))
print('All unique Clickstream Data Actions:', unique_clickstream_data_actions)

```

All unique Service Interaction Types: ['Chat', 'Call', 'Email']
All unique Payment History Methods: ['PayPal', 'Bank Transfer', 'Credit Card']
All unique Clickstream Data Actions: ['Add to Cart', 'Search', 'Click']

FROM THESE THREE COLUMNS, I WILL FUTHER ENCODE MORE

```

In [15]: #ServiceInteractions
for usit in unique_service_interaction_type:
    df[f'ServiceInteractions_{usit}'] = df['ServiceInteractions'].apply(lambda x: sum(1 for i in x if i['InteractionType'] == usit))

# Payment History
df['PaymentHistoryNoOfPayments'] = df['PaymentHistory'].apply(lambda x: sum(i['NoOfPayments'] for i in x))
df['PaymentHistoryAvgNoOfLatePayments'] = df['PaymentHistory'].apply(lambda x: np.mean([i['AvgNoOfLatePayments'] for i in x]))

#ClickStreamData
for ucda in unique_clickstream_data_actions:
    df[f'ClickStreamData_{ucda}'] = df['ClickstreamData'].apply(lambda x: len([i for i in x if i['Action'] == ucda]))

```

See a subset of the new data

```
In [16]: df.head()
```

Out[16]:

	CustomerID	Name	Age	Gender	Location	Email
0	1001	Mark Barrett	31	Male	Andrewfort	allison74@example.net
1	1002	Jeremy Welch	66	Female	Millerhaven	fmiller@example.com
2	1003	Brandon Patel	36	Female	Lozanostad	jasonbrown@example.org
3	1004	Tina Martin	62	Female	South Dustin	matthew62@example.net
4	1005	Christopher Rodriguez	68	Female	West James	shannonstrickland@example.org

5 rows × 46 columns



See all columns

In [17]: `df.columns`

Out[17]: Index(['CustomerID', 'Name', 'Age', 'Gender', 'Location', 'Email', 'Phone',
 'Address', 'Segment', 'PurchaseHistory', 'SubscriptionDetails',
 'ServiceInteractions', 'PaymentHistory', 'WebsiteUsage',
 'ClickstreamData', 'EngagementMetrics', 'Feedback',
 'MarketingCommunication', 'NPS', 'ChurnLabel', 'Timestamp', 'MonthYear',
 'FeedbackRating', 'PurchasProducts', 'PurchaseFrequency',
 'PurchaseValue', 'SubscriptionPlan', 'SubscriptionStartDate',
 'SubscriptionEndDate', 'SubscriptionDuration', 'WebsitePageViews',
 'WebsiteTimeSpent', 'EngagementMetricsLogins',
 'EngagementMetricsFrequency', 'FeedbackComment',
 'MarketingCommunicationNoOfEmails',
 'MarketingCommunicationOpenClickDiff',
 'MarketingCommunicationSentOpenDiff', 'ServiceInteractions_Chat',
 'ServiceInteractions_Call', 'ServiceInteractions_Email',
 'PaymentHistoryNoOfPayments', 'PaymentHistoryAvgNoOfLatePayments',
 'ClickStreamData_Add to Cart', 'ClickStreamData_Search',
 'ClickStreamData_Click'],
 dtype='object')

Pick out some Columns

In [18]: `df_ = df[[
 'Age',`

```

'Gender',
'NPS',
'ChurnLabel',
'PurchaseFrequency',
'PurchaseValue',
'SubscriptionPlan',
'WebsitePageViews',
'WebsiteTimeSpent',
'EngagementMetricsLogins',
'EngagementMetricsFrequency',
'FeedbackRating',
'MarketingCommunicationNoOfEmails',
'MarketingCommunicationOpenClickDiff',
'MarketingCommunicationSentOpenDiff',
'ServiceInteractions_Email',
'ServiceInteractions_Chat',
'ServiceInteractions_Call',
'PaymentHistoryNoOfPayments',
'ClickStreamData_Search',
'ClickStreamData_Click',
'ClickStreamData_Add to Cart',
'SubscriptionDuration'
]]
df_.head()

```

Out[18]:

	Age	Gender	NPS	ChurnLabel	PurchaseFrequency	PurchaseValue	SubscriptionPlan
0	31	Male	3	1	38	3994.72	Express
1	66	Female	6	0	4	2844.35	Pro
2	36	Female	3	0	14	1866.52	Essential
3	62	Female	1	1	28	1378.64	Smart
4	68	Female	3	0	39	2425.05	Basic

5 rows × 23 columns



Let's see the name of the column we have now

In [19]: df_.columns

Out[19]:

```

Index(['Age', 'Gender', 'NPS', 'ChurnLabel', 'PurchaseFrequency',
       'PurchaseValue', 'SubscriptionPlan', 'WebsitePageViews',
       'WebsiteTimeSpent', 'EngagementMetricsLogins',
       'EngagementMetricsFrequency', 'FeedbackRating',
       'MarketingCommunicationNoOfEmails',
       'MarketingCommunicationOpenClickDiff',
       'MarketingCommunicationSentOpenDiff', 'ServiceInteractions_Email',
       'ServiceInteractions_Chat', 'ServiceInteractions_Call',
       'PaymentHistoryNoOfPayments', 'ClickStreamData_Search',
       'ClickStreamData_Click', 'ClickStreamData_Add to Cart',
       'SubscriptionDuration'],
      dtype='object')

```

```
In [20]: df_.loc[0] #this shows if all column values are in numeric numbers suitable for
```

```
Out[20]: Age                      31
Gender                   Male
NPS                       3
ChurnLabel                 1
PurchaseFrequency          38
PurchaseValue            3994.72
SubscriptionPlan        Express
WebsitePageViews           49
WebsiteTimeSpent            15
EngagementMetricsLogins      19
EngagementMetricsFrequency   Weekly
FeedbackRating                  1
MarketingCommunicationNoOfEmails    8
MarketingCommunicationOpenClickDiff  319.0
MarketingCommunicationSentOpenDiff  818.0
ServiceInteractions_Email       1
ServiceInteractions_Chat        2
ServiceInteractions_Call        1
PaymentHistoryNoOfPayments      40
ClickStreamData_Search          12
ClickStreamData_Click             4
ClickStreamData_Add to Cart      8
SubscriptionDuration          871
Name: 0, dtype: object
```

We need to encode all variables in strings to numerical value to be able to use our ML prediction metrics. Also to check if we have excess variables to be encoded.

Now Checking for unique values in each column

```
In [21]: print('Total dataset length', len(df_))

df_[['Gender', 'SubscriptionPlan', 'EngagementMetricsFrequency']].nunique()
```

```
Total dataset length 12483
```

```
Out[21]: Gender                  2
SubscriptionPlan                20
EngagementMetricsFrequency       3
dtype: int64
```

Encoding the string parameters.

```
In [22]: #Gender encoding
gender_map = {'Male': 0, 'Female': 1}

#Subscription encoding
unique_subscription_plans = df_['SubscriptionPlan'].unique()
subscription_plan_map = {unique_subscription_plans[i]: i for i in range(len(unique_subscription_plans))}

#EngagementMetricsFrequency
unique_engagement_frequency = df_['EngagementMetricsFrequency'].unique()
engagement_frequency_map = {unique_engagement_frequency[i]: i for i in range(len(unique_engagement_frequency))}

#Encode
df_.loc[:, 'Gender'] = df_.loc[:, 'Gender'].map(gender_map)
```

```
df_.loc[:, 'SubscriptionPlan'] = df_.loc[:, 'SubscriptionPlan'].map(subscription_map)
df_.loc[:, 'EngagementMetricsFrequency'] = df_.loc[:, 'EngagementMetricsFrequency'].map(engagement_frequency_map)
```

In [23]: `df_.loc[0]`


Out[23]:

Age	31
Gender	0
NPS	3
ChurnLabel	1
PurchaseFrequency	38
PurchaseValue	3994.72
SubscriptionPlan	0
WebsitePageViews	49
WebsiteTimeSpent	15
EngagementMetricsLogins	19
EngagementMetricsFrequency	0
FeedbackRating	1
MarketingCommunicationNoOfEmails	8
MarketingCommunicationOpenClickDiff	319.0
MarketingCommunicationSentOpenDiff	818.0
ServiceInteractions_Email	1
ServiceInteractions_Chat	2
ServiceInteractions_Call	1
PaymentHistoryNoOfPayments	40
ClickStreamData_Search	12
ClickStreamData_Click	4
ClickStreamData_Add to Cart	8
SubscriptionDuration	871
Name:	0, dtype: object

Plot correlation Matrix

In [24]:

```
df_corr = df_.corr()
fig, ax = plt.subplots(figsize=(13, 13))
sns.heatmap(df_corr, annot=True, fmt='.1f', ax=ax)
plt.show()
```



Split data into train, test and validation sets.

```
In [25]: X = df_.drop(columns=['ChurnLabel'])
y = df_['ChurnLabel']

X_train, X_other, y_train, y_other = train_test_split(X, y, train_size=0.8, random_state=42)
X_test, X_val, y_test, y_val = train_test_split(X_other, y_other, test_size=0.3, random_state=42)
```

```
In [29]: ss = StandardScaler()
X_train = ss.fit_transform(X_train)
X_val = ss.transform(X_val)
X_test = ss.transform(X_test)
```

Modelling

Two different models for modelling:

- LogisticRegression,
- DecisionTreeClassifier.

Metrics:

- Accuracy score,
- Precision score,
- Recall score,
- F1 score.

```
In [34]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

def evaluate(X, y, model, subset=''):
    y_pred = model.predict(X)

    print(f'{subset} Accuracy Score: {accuracy_score(y, y_pred)}')
    print(f'{subset} Precision Score: {precision_score(y, y_pred)}')
    print(f'{subset} Recall Score: {recall_score(y, y_pred)}')
    print(f'{subset} F1 Score: {f1_score(y, y_pred)}')
```

Modelling with Logistic Regression

```
In [38]: #Build the model
lr = LogisticRegression()
lr.fit(X_train, y_train)

# Evaluate the model on train and validation subsets
evaluate(X_train, y_train, lr, 'Train')
evaluate(X_val, y_val, lr, 'Validation')
```

```
Train Accuracy Score: 0.9709593430803124
Train Precision Score: 0.9767071471232331
Train Recall Score: 0.965938176806458
Train F1 Score: 0.9712928133042962
Validation Accuracy Score: 0.968
Validation Precision Score: 0.9643835616438357
Validation Recall Score: 0.9696969696969697
Validation F1 Score: 0.967032967032967
```

Modelling with Decision Tree

```
In [46]: #Build the model
dt = DecisionTreeClassifier(max_depth=5)
dt.fit(X_train, y_train)
/

# Evaluate on the train and validation subsets
evaluate(X_train, y_train, dt, subset='Train')
evaluate(X_val, y_val, dt, subset='Validation')
```

```
Train Accuracy Score: 0.9768676146605247
Train Precision Score: 0.9775413711583925
Train Recall Score: 0.9769639692852924
Train F1 Score: 0.9772525849335303
Validation Accuracy Score: 0.968
Validation Precision Score: 0.9643835616438357
Validation Recall Score: 0.9696969696969697
Validation F1 Score: 0.967032967032967
```

Evaluation on the Test Set

```
In [47]: evaluate(X_test, y_test, lr, 'LogisticsRegression Test')
evaluate(X_test, y_test, dt, 'DecisionTreeClassifier Test')
```


```
LogisticsRegression Test Accuracy Score: 0.9690898683457355
LogisticsRegression Test Precision Score: 0.9733796296296297
LogisticsRegression Test Recall Score: 0.9644495412844036
LogisticsRegression Test F1 Score: 0.9688940092165899
DecisionTreeClassifier Test Accuracy Score: 0.9725243274184316
DecisionTreeClassifier Test Precision Score: 0.9735632183908046
DecisionTreeClassifier Test Recall Score: 0.9713302752293578
DecisionTreeClassifier Test F1 Score: 0.9724454649827784
```

Plotting Confusion Matrix

```
In [49]: lr_y_pred = lr.predict(X_test)
logistic_regression_confusion_matrix = confusion_matrix(y_test, lr_y_pred)

display = ConfusionMatrixDisplay(confusion_matrix=logistic_regression_confusion_
display.plot()
plt.title('Confusion Matrix on Test Set for Logistic Rrgression')
plt.show()
```


Confusion Matrix on Test Set for Logistic Rrgression



```
In [51]: dt_y_pred = dt.predict(X_test)
decision_tree_confusion_matrix = confusion_matrix(y_test, dt_y_pred)

display = ConfusionMatrixDisplay(confusion_matrix=decision_tree_confusion_matrix)
display.plot()
plt.title('Confusion Matrix on Test Set for Decision Tree')
plt.show()
```

Confusion Matrix on Test Set for Decision Tree



Conclusion

The most important Features the affects and correlates Churn Label are:

- The number of service interaction the customer had through Call Email, Chat,
- The number of times the customers has made last payments,
- The time customer spends on the company websites,
- The NET Promoter Score (NPS).

In []: